

Maincode Homework

Task

Dr Brian Le

December 10, 2025

Contents

| | |
|--|----|
| 1. Task Description | 3 |
| 2. Dataset statistics | 4 |
| 3. Data Pipeline | 6 |
| 3.1. Data Cleaning and Normalisation | 6 |
| 3.1.1. Deduplication | 7 |
| 3.1.2. Normalisation | 7 |
| 3.1.3. Text Filtering | 7 |
| 3.1.4. Language Detection | 8 |
| 3.1.5. PII Removal | 9 |
| 3.2. Data Preparation for Training | 9 |
| 4. Results | 10 |
| 4.1. Analysis and Visualisation | 10 |
| 4.1.1. Language Detection Results | 11 |
| 4.1.2. Text Length Distribution | 11 |
| 4.1.3. Word Count Distribution | 11 |
| 4.1.4. PII Hit Rates | 11 |
| 5. Future work | 12 |
| Bibliography | 13 |
| Index of Tables | 14 |

1. Task Description

The task given was to build an end-to-end data pipeline from a unprepared dataset. The task specifies the final dataset must be an English language dataset with the following criteria:

1. Data acquisition
2. Data cleaning, normalisation and tokenisation
3. Training ready exports (i.e. tokenised, mixtures, shards, etc.)
4. Inspectability (e.g. histograms of length, lang scores, dup markers, PII hit-rates, drop reasons)
5. Conceptual plan for scaling

The methodology and implementation of the data pipeline will be described in this report. Where applicable, decision points will be outlined along with the reasoning behind them.

2. Dataset statistics

The dataset provided consisted of 269378 rows of `jsonl` data with two columns (`text` and `url`).

From a cursory glance, the dataset consists of a mixture of data from various sources. The URL field contains 34% of null values while the text field is complete. A total of 63599 sources are represented in the dataset with the top 20 sources being:

| Domain | Rows |
|----------------------------|-------|
| No source | 90519 |
| github.com | 71415 |
| www.taiwannews.com.tw | 3030 |
| en.wikipedia.org | 2326 |
| www.nigeriatoday.ng | 2162 |
| placeholder.co | 2110 |
| abcnews.go.com | 2060 |
| sample-company.net | 1997 |
| example.com | 1991 |
| testsite.org | 1964 |
| demo-page.info | 1938 |
| www.yahoo.com | 946 |
| uproxx.com | 844 |
| www.israelnationalnews.com | 746 |
| www.nytimes.com | 735 |
| www.wafb.com | 650 |
| newyork.cbslocal.com | 478 |
| www.ocregister.com | 475 |
| www.engadget.com | 452 |
| www.parrysound.com | 416 |

Table 1: Top 20 domains

The presence of placeholder domains such as `placeholder.co`, `example.com`, and `testsite.org` indicates that some data entries may be synthetic or used for testing purposes. This will be used as a metric to determine the efficacy of the data cleaning process. Sources such as `github.com` are likely to contain code snippets or technical documentation so may need to be treated differently during preprocessing.

Of the non-null URLs, the top 20 top level domains (TLDs) are:

| Domain | Rows |
|--------|--------|
| com | 128768 |
| org | 10330 |
| de | 6303 |
| net | 4722 |
| uk | 3924 |
| tw | 3056 |
| info | 2764 |
| co | 2503 |
| ng | 2203 |
| au | 1567 |
| es | 1387 |
| ca | 1210 |
| edu | 822 |
| at | 559 |
| eu | 513 |
| mx | 456 |
| ch | 444 |
| in | 411 |
| sg | 394 |
| ru | 380 |

Table 2: Top 20 top level domains

This indicates an admixture of English (com , org , net , uk , au , ca , edu , eu , mx , in , sg), German (de , at , ch), Spanish (es) and other languages (tw , ng , ru). The cleaning is likely to involve language detection and filtering to ensure the final dataset is exclusively English.

3. Data Pipeline

The data pipeline will be divided into two stages:

1. Data Cleaning and Normalisation
2. Data Preparation for Training
3. Analysis and Visualisation of Results

For the pipeline implementation, `python3.13` will be used along with the following libraries:

- `pandas` : for data ingestion, manipulation and analysis
- `langdetect` : for language detection
- `lingua` : for language detection
- `BeautifulSoup` : for HTML tag removal
- `nltk` : for tokenisation and text preprocessing
- `presidio` : for PII detection and removal
- `matplotlib` and `seaborn` : for data visualisation and inspection

Time did not allow for experimentation of `detoxify` for removal of toxic content but this would be a useful addition to the cleaning pipeline. The laptop I was working on has poor support for the `pytorch` version needed for `detoxify` so I was unable to get it working in time.

Given the small size of the dataset, the entire dataset will be processed in-memory using `pandas` `DataFrames`. The `pipe` functionality of `pandas` will be used to chain together the various cleaning steps in a modular fashion.

Ideally this would be scaled using a distributed data processing framework such as `Apache Spark`, `Ray` or `Dask` for larger datasets.

A simple `json` based configuration file will be used to specify the parameters for each cleaning step (as well as which steps to run).

3.1. Data Cleaning and Normalisation

The pipeline aims to filter and clean the dataset to meet the following criteria:

- Exclusively English language Text
- Removal of PII
- Deduplicated text
- Removal of non-alphabetic text
- Normalised text

The ordering of the cleaning steps is important to ensure minimising runtime costs. The bottleneck for this was the PII detection, the language detection and the cleaning of the text inputs. Each step which filters out rows has a flag for whether to apply the filter after calculating.

The cleaning steps are as follows:

1. `deduplicate_data`

2. clean_text_ascii
3. check_text_is_alphabetic
4. check_text_is_hyperlink
5. clean_text_html
6. clean_short_length
7. extract_domain_from_col
8. detect_pii (not by default)

3.1.1. Deduplication

The first step in the cleaning pipeline is deduplication of the dataset based on the `text` field. Only exact duplicates are removed at this stage to reduce the dataset size for subsequent processing steps.

A native `pandas` function is used for this:

```
df.duplicated(subset=["text"])
```

Potential future work could involve fuzzy matching or semantic similarity to identify near-duplicates. This would become more computationally expensive so was not implemented in this initial pipeline.

This step needs to be run first to reduce the dataset size for subsequent processing steps and can't be batched as with other steps.

3.1.2. Normalisation

As this is a natural language dataset, usual normalisation steps of lowercasing, punctuation removal, and whitespace trimming would potentially prevent training of a robust language model. Therefore, only leading and trailing whitespace is removed from the text entries.

Further cleaning focussed on removal of HTML tags and non-printable characters.

The `BeautifulSoup` library is used for HTML tag removal as it is robust and can handle malformed HTML. Note this step should also consider XML tags if the dataset contains XML data. Given it appears to have a very low amount of XML data, this was not implemented in this initial pipeline.

For non-printable character removal, a native `python` function is used to filter out characters outside the printable ASCII range.

3.1.3. Text Filtering

Several steps are implemented to filter out unwanted text entries that are not helpful for training a language model.

1. Non-alphabetic text removal: Entries that contain a high proportion of non-alphabetic characters (e.g. numbers, symbols) are removed. A threshold of 70% alphabetic characters is used.

2. Hyperlink removal: Entries that are primarily hyperlinks or contain a high proportion of URLs are removed. A threshold of 50% URL content is used.
3. Short length removal: Entries that are too short (less than 20 characters) are removed as they are unlikely to provide meaningful context for training.
4. Removal of code snippets: Given the presence of `github.com` as a major source, code snippets may be present in the dataset

Some examples of filtered text entries:

```
2016
+1
http://daytonflyers.com/roster.aspx?rp_id=2579
+1
https://www.youtube.com/watch?v=9n1cqbvCpDg&ebc=ANyPxKrVp5Vbw0bmXfouF74NnPJYew2qcz9QyWDyxk2
FTFJT1Xfc28qai251XUg0A8XqIRJ5EMNOhyT09L1s3YaS1SpaZHHTA
;)
```

Non-alphabetic text is removed by checking whether any character in the text is an alphabetic character using the native `str.isalpha()` function. A proportion approach (eg. > 80%) is probably preferred but for brevity this was not implemented.

Several texts are predominantly composed of a single (or multiple) hyperlink(s). These are identified using a regex pattern (`re.findall(r"(?P<url>https?:\/\/[^s]+)", text)`) to match URLs and calculating the proportion of URL content in the text.

Short length texts are removed using a simple length check on the text string. This length threshold is calculated using both the text length and the number of words (split by whitespace). The thresholds used were chosen rather arbitrarily and could be tuned further.

Finally to remove code snippets, `github.com` sourced texts are filtered out. This represents a blunt approach and could be improved by implementing code detection using regex patterns or language models trained to identify code snippets. The main concern is the other steps for cleaning may inadvertently modify code snippets which make them more difficult to use in training. Alternatively this could be simply flagged and then treated separately during training. Further work could involve identifying the programming language used with the `pygments` library.

3.1.4. Language Detection

To ensure the dataset is exclusively English language text, language detection is performed using two libraries: `langdetect` and `lingua`.

The `langdetect` library is used first as it is lightweight and provides a very simple implementation that can return multiple language matches with associated probabilities. This library can not be vectorised so the row-by-row calculation makes this a very slow step in the pipeline. As this library has little fault tolerance for noisy text, this package was used to filter out obvious non-English text for further study.

The `lingua` library is then used as an alternative option as it provides higher accuracy at the cost of increased computational resources. The `lingua` library also allows for detecting from a subset of languages and has an interface for batch processing. `lingua` also contains a quicker detection mode which sacrifices some accuracy for speed.

Another library `fast-langdetect` ^o claims to provide high speed language detection using `fasttext` models but was not explored in this initial pipeline.

A final library `langid` was considered but not explored due to time constraints. The `langid` library contains a larger set of languages (97) compared to `langdetect` (55) and `lingua` (75) so may provide better coverage for edge cases. For the sake of determining only English language text, this was not explored further.

The decision was to use `lingua` for the final language detection step as it provided the best balance of accuracy and speed for this dataset.

3.1.5. PII Removal

Personally Identifiable Information (PII) removal is a step commonly performed to cleanse a dataset of sensitive information prior to training. The PII library used is the `presidio` library from Microsoft. This library is widely implemented including for Azure OpenAI services. It utilises a mixture of named entity recognition (NER) models, regex and rule based logic to identify and redact PII from text.

This step is the most computationally expensive step in the pipeline so is performed after deduplication and language filtering to reduce the number of rows processed.

An important consideration is whether to remove rows containing PII or to redact the PII from the text. As there is a risk of losing too much data, the decision was made to redact the PII from the text rather than remove the entire row. There is a risk that if there are too many names in a given text, the context may be lost as multiple redactions result in incoherent text (eg. “John went to the store. He met with Mary and they talked about Peter.” becomes “<PERSON> went to the store. He met with <PERSON> and they talked about <PERSON> .”). Note this step is not enabled by default due to the computational cost.

3.2. Data Preparation for Training

To prepare for training, the dataset is shuffled and tokenised. Shuffling is performed using the native `pandas` function:

```
df.sample(frac=1, random_state=42)
```

If particular admixtures need to be maintained (eg. domain specific data), stratified sampling could be implemented here.

4. Results

After running the cleaning pipeline with the default settings (deduplication, normalisation, text filtering, language detection), the final dataset statistics are as follows:

| Metric | Value |
|---------------------|--------|
| Initial Rows | 269378 |
| Final Rows | 123456 |
| Rows Removed | 145922 |
| Percentage Retained | 45.8% |

Table 3: Final dataset statistics

The final dataset consists of 123456 rows of English language text.

4.1. Analysis and Visualisation

To inspect the results of the cleaning pipeline, several visualisations were created using `matplotlib` and `seaborn`.

The final top 20 domains are as follows:

| Metric | Value |
|---------------------|--------|
| Initial Rows | 269378 |
| Final Rows | 123456 |
| Rows Removed | 145922 |
| Percentage Retained | 45.8% |

Table 4: Top 20 domains after cleaning

From the results we see no further presence of placeholder domains. The remaining top 10 TLDs are:

| Domain | Rows |
|--------|-------|
| .com | 47267 |
| .org | 7506 |
| .uk | 3867 |
| .tw | 2957 |
| .net | 2016 |
| .ng | 1766 |
| .au | 1512 |
| .ca | 1019 |

| Domain | Rows |
|--------|------|
| edu | 768 |
| info | 564 |

Table 5: Top 10 top level domains after cleaning

4.1.1. Language Detection Results

A bar chart showing the distribution of detected languages before and after cleaning was created to visualise the effectiveness of the language filtering step.

4.1.2. Text Length Distribution

A histogram of text lengths (in characters) before and after cleaning was created to visualise the impact of the cleaning steps on text length.

4.1.3. Word Count Distribution

4.1.4. PII Hit Rates

5. Future work

The primary concern for this pipeline is scaling to larger datasets. The current implementation processes the entire dataset in-memory using `pandas` DataFrames.

For larger datasets some batch processing or distributed data processing framework would be needed. Options include:

- Apache Spark : widely used distributed data processing framework with support for Python via `PySpark`.
- Dask : parallel computing library that integrates well with `pandas` and can scale from a single machine to a cluster.
- Ray : distributed computing framework that can be used for parallel data processing
- Prefect or Airflow : workflow orchestration tools that can manage complex data pipelines and integrate with distributed processing frameworks

Of note the particular steps which need efficient batching are language detection and PII removal.

The `detoxify` library for toxic content removal would also be a useful addition to the cleaning pipeline. Some work is required to get this working with my particular setup but should in principle be straightforward to implement.

Bibliography

Index of Tables

| | |
|---|----|
| Table 1 Top 20 domains | 4 |
| Table 2 Top 20 top level domains | 5 |
| Table 3 Final dataset statistics | 10 |
| Table 4 Top 20 domains after cleaning | 10 |
| Table 5 Top 10 top level domains after cleaning | 10 |